

Time

System calls for Time

`stime` allows the **superuser** to set a global kernel variable to a value that gives the current time.

```
#include <time.h>
stime(pvalue);
```

where `pvalue` **points to a long integer** that gives the time measured in seconds from midnight before (00:00:00) January 1, 1970, GMT. The clock interrupt handler increments the kernel variable once a second.

Time retrieves the time as set by `stime()`:

```
#include <time.h>
time(tloc);
```

where `tloc` points to a location in the user process for return value.

`times` retrieves the **cumulative times that the calling process spent executing in kernel mode and user mode** and the cumulative times that all zombie children had executed in user mode and kernel mode. `times` **returns the elapsed time** "from an arbitrary point in the past", **usually the time of system boot**.

```
#include <sys/times.h>
times(tbuffer);
struct tms *tbuffer;

struct tms{
/* time_t is the data structure for time */
time_t tms_utime; /* user time of process */
time_t tms_stime; /* kernel time of process */
time_t tms_cutime; /* user time for children */
time_t tms_cstime; /* kernel time for children */
}
```

The child times **do not include time spent in the `fork()` and `exit()`**, and all **times can be distorted by times spent handling interrupts or doing context switches.**

Alarm system call set the alarm clock of a process. User processes can schedule **alarm signals** using the alarm system calls:

```
#include <unistd.h>
int alarm(int sec);
```

generates a SIGALRM signal for the process after the number of real-time seconds specified by `sec` (seconds).

If there is a previous `alarm()` request with time remaining, `alarm()` returns non zero that is the number of seconds until the previous request would have generated a SIGALRM signal. Otherwise, `alarm()` returns 0.

Example:

```
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/signal.h>

main(argc, argv)
    int argc;
    char *argv[];
{
    extern unsigned alarm();
    extern wakeup();
    struct stat statbuf;
    time_t axtime;

    if (argc != 2)
    {
        printf("only 1 arg\n");
        exit();
    }

    axtime = (time_t) 0;
    for (;;)
    {
        /* find out file access time */
```

```

        if (stat(argv[1], &statbuf) == -1)
        {
            printf("file %s not there\n", argv[1]);
            exit();
        }
        if (axtime != statbuf.st_atime)
        {
            printf("file %s accessed\n", argv[1]);
            axtime = statbuf.st_atime;
        }
        signal(SIGALRM, wakeup);      /* reset for alarm */
        alarm(60);
        pause();                      /* sleep until signal */
    }
}

wakeup()
{
}

```

Clock

The functions of the **clock handler** are to:

- restart the clock
- schedule invocation of internal kernel functions based on internal timers
- provide execution profiling capability for the kernel and for user process
- gather system and process accounting statistics
- keep track of time
- send alarm signals to processes on request
- periodically wake up the swapper process
- control process scheduling

Algorithm for clock handler

```

algorithm clock
input: none
output: none
{
    restart clock;
}

```

```

    if(callout table not empty){
        adjust callout times;
        schedule callout function if time elapsed;
    }
    if(kernel profiling on)
        note program counter at time of interrupt;
    if(user profiling on)
        note program counter at time of interrupt;
    gather system statistics;
    gather statistics per process;
    adjust measure of process CPU utilization;
    if(1 second or more since last here and interrupt not in critical region of
code){
        for(all processes in the system){
            adjust alarm time if active;
            adjust measure of CPU utilization;
            if(process to execute in user mode)
                adjust process priority;
        }
        wakeup swapper process is necessary;
    }
}

```

Restarting the clock

- When the clock interrupts the system, most machines require that the clock be reprimed (reprepare) by software instructions so that it will interrupt the processor again after a suitable interval.
- Such instructions are hardware dependent.

Internal System Timeouts

- Some kernel operations, particularly device drivers and network protocols, require the invocation of kernel functions on a real time basis.
- Example: Putting kernel into raw mode so that the kernel satisfies user read requests at fixed intervals.
- The kernel stores the necessary information in the **callout table**, which consists of the **function to be invoked** when the time expires, a **parameter for the function**, and the **time** in clock ticks **until the function should be called**. User has no direct control over this.

Function	Time to Fire
a()	-2
b()	3
c()	10

Before

Function	Time to Fire
a()	-2
b()	3
f()	2
c()	8

After

Figure 8.10. Callout Table and New Entry for f

ADDING A NEW ENTRY

- The kernel finds the correct (timed) position to insert the new entry and **appropriately adjusts the time field of the entry immediately after the new entry**.
- New entry: *f* with time 5s. Now *c()* will execute in 13s ($10 + 3$)
- It finds that it needs to execute after *b()* to satisfy the 5s time ($3 < 5$), its time will be now

$$5 - 3 = 2$$

and the entry immediately after

$$10 - 2 = 8$$

so that *c()* will still fire in 13s.

- So the time for the new entry is time - time for previous entry.

DATA STRUCTURES FOR THIS

- Linked list can be used
- Or the kernel can read just the position of the entries when changing the table (not too expensive if the kernel does not use the callout table too much).

Procedure

- At every clock interrupt, the clock handler checks if there are any entries in the callout table
- If there are any, it decrements the time field of the first entry.
- Because of the way, the kernel keeps the entries, decrementing the first entry effectively decrements the time field for all entries in the table.

- If the time field of the first entry < 0 , the specified function should be invoked.
- The clock handler does not directly invoke the function since it will block later clock interrupts, instead it schedules the function by generating a "software interrupt"
- Because, s/w interrupts are at lower priority than other interrupts, they are blocked until kernel handles all other interrupts.
- The entry is then removed.

Profiling

- Gives a measure of how much time the system is executing in user mode vs. kernel mode, and how much time it spends executing individual routines in the kernel.
- The kernel profile driver monitors the relative performance of kernel modules by sampling system activity at the time of the clock interrupt.
- The profile driver has a list of kernel address to sample, usually addresses of kernel function.
- If the kernel profiling is enabled, **the clock handler invokes the interrupt handler of the profile driver**, which determines the whether the processor mode at the time of interrupt was user or kernel.
- If the mode was user, the **profiler increments a count for user execution**, but if the mode was kernel, it **increments an internal counter corresponding to the program counter**.
- User processes can read the profile driver to obtain kernel counts and do statistical measurements.

Clock handler invokes interrupt handler of profile driver → determines processor mode (user or kernel) → increments a counter

Profile execution of processes at user-level with `profil()`

```
profil(buff, bufsz, offset, scale);
```

where `buff` is the address of an array in user space and `bufsz` is its size

`offset` is the virtual address of user subroutine

`scale` is a factor that maps user virtual addresses into array.

In the user mode, the clock handler examines the user program counter at the time of interrupt, compares it to `offset`, and increments a location in `buff` whose address is a function of `bufsz` and `scale`.

compare `offset` and `pc` → increments a location in `buff`

Algorithm	Address	Count
bread	100	5
breada	150	0
bwrite	200	0
brelease	300	2
getblk	400	1
user	—	2

Figure 8.11. Sample Addresses of Kernel Algorithms

Accounting and Statistics

- Every process has 2 fields in its `u` area to **keep a record of elapsed kernel** and **user time**.
- When handling clock interrupts, the kernel updates the appropriate field for the executing process, depending on whether it was executing in the user mode or the kernel mode.
- Parent processes gather statistics for their child processes in the `wait()`.
- Every process has one field in its `u` area for **the kernel to log its memory usage**.
- When the clock interrupts, the kernel calculates the total memory used by a process as a function of its private memory regions and its proportional usage of the shared memory regions (since a process may use a part of the shared memory region).
- Example: shared region of size = $50K$ with 4 other processes and text and data regions of size $25K$ and $40K$ respectively. So memory usage equals,

$$memory = 50K/5 + 25K + 40K = 75K$$